



Ansible in 4 Hours  
Sander van Vugt  
[mail@sandervanvugt.nl](mailto:mail@sandervanvugt.nl)

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles, all rendered in a light gray color.

# Ansible in 4 Hours

Course information

# Your Instructor

- Sander van Vugt
- [mail@sandervanvugt.nl](mailto:mail@sandervanvugt.nl)
- <http://sandervanvugt.com>
- Author of
  - Beginning Ansible
  - Automation with Ansible (expected December 2018)
  - Ansible Certification (expected December 2018)
- Available as recorded video courses on Safaribooksonline
- Also attend my Safaribooksonline Live session "Automation with Ansible"

# Agenda

- What is Ansible?
- Installing Ansible
- Configuring Ansible Managed Servers
- Running Ansible Ad-hoc Commands
- Running Ansible Playbooks
- Using Variables
- Using Inclusions

# Course Files

- The demo files used in this course are available on github
- Use **git clone <https://github.com/sandervanvugt/ansible-3h>** to download

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles, all rendered in a light gray color.

# Ansible in 4 Hours

## Understanding Ansible

# What is Ansible?

- Ansible is a Configuration Management System, written in Python
- It uses playbooks to run tasks in a way that always gives the same result
- It is push-based
- It doesn't require any agent on the managed node
- Apart from managing computers, it's doing very well in managing network devices as well
- Configurations are written in YAML
- Ansible generated Python scripts that are started on the managed nodes

# Ansible compared to others

- YAML is easy to learn and read
- No agent to install on managed hosts
  - You'll need Python and SSH though
- Push based, which gives you more control over the process
  - An optional ansible-pull tool is available for if you want to be able to pull configurations
- Many modules are available
- Idempotent: running the same playbook multiple times will give you the same results



# Required Skills

- Use SSH
- Use Linux Commands
- Install software
- Use sudo
- Manage permissions
- Manage services
- Work with variables

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles, with the outermost circle being the most prominent.

# Ansible in 4 Hours

## Installing Ansible

# Installing the Ansible Controller

- On Linux: use the version in the repositories
  - Might bring you a somewhat older version
  - Easiest and for that reason recommended
- On MacOS, use Homebrew package manager
- Or else, use the Python package manager pip

# Configuring SSH to Manage Hosts

- Set up SSH Key-based authentication
  - **ssh-keygen**
- This creates a public key as well as a private key
  - The server that has the public key sends a challenge that can only be answered with the private key
  - Keep the private key in the local user account on the control node
  - Send the public key to the `~/.ssh/authorized_keys` file in the target user home directory
    - Use **ssh-copy-id user@remotehost**
    - Notice that the local user name and the remote user name do NOT have to be the same

# The Inventory File

- The inventory file identifies and groups managed hosts
- The inventory file may be indicated with the `-i` option
- Typically, you would create an Ansible project directory in your home directory, and put an inventory file in there
- You can specify which inventory to use in the `ansible.cfg` file
- For use in more advanced projects, multiple inventory files may be used

# Managing Managed Hosts

- After installation, you can use the **ansible** command to run ad-hoc commands against remote hosts
- Remote hosts need to be specified in the inventory file
- The inventory file allows you to define managed hosts
- Hosts are specified by their name or IP address
- Hosts may be mentioned in the inventory more than once
  - This allows you to create logical groups
- In **ansible** commands, you'll mention host names, as well as the inventory file that you're going to use
  - **ansible server1.example.com,server2.example.com -i myinventory --list-hosts**

# Lab 1: Installing Ansible

1. **useradd ansible; passwd ansible; su - ansible**
2. On both nodes: **sudo yum install python2 epel-release -y**
3. Remaining steps on control: **sudo yum install -y ansible**
4. **ssh server1.ansible.local**
5. **ssh-keygen**
6. **ssh-copy-id server1.ansible.local**
7. **mkdir ~/install**
8. **vim ~/install/inventory**

control.ansible.local

server1.ansible.local

9. **ansible all -i inventory --list-hosts**

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

# Ansible in 4 Hours

## Using Inventory



# Understanding Inventory

- Ansible uses an inventory file, which must be used to identify managed hosts
- The location of the inventory file can be anywhere and is specified in the `ansible.cfg` file
  - `/etc/ansible/hosts`
  - current project directory
  - specified with the `-i` option while running Ansible commands
- Inventory files may be statically created or dynamically generated
  - Static inventory works for small environments
  - Dynamic inventory uses 3<sup>rd</sup> party scripts to identify hosts in a specific environment

# Working with Dynamic inventory

- When using the **ansible** command, use the **-i** option, followed by the inventory script you'd like to use
  - Ensure that the inventory script is executable
- Write your own script or use a script that is available for the different externally supported cloud environments

# Using Groups in Inventory

- An inventory file contains a list of hosts
- Hosts may be grouped to make referring to hosts easier
- A host can be a part of multiple groups
- The host group **all** is always present and doesn't have to be defined

# Nesting Host Groups

- Host Groups may be nested in inventory

```
[webservers]  
web1.example.com  
web2.example.com
```

```
[dbservers]  
db1.example.com  
db2.example.com
```

```
[servers:children]  
webservers  
dbservers
```

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

# Ansible in 4 Hours

Using `ansible.cfg`

# The ansible.cfg file

- The ansible.cfg file is used to specify generic settings
  - How to escalate permissions
  - Where to find the inventory file
  - And more
- The following locations are used
  - `$ANSIBLE_CONFIG`
  - `./ansible.cfg`
  - `~/.ansible.cfg`
  - `/etc/ansible/ansible.cfg`
- It is common practice to put it in the current project directory
- Using section headers is important!

# Common ansible.cfg parameters

```
[defaults]  
inventory = /etc/ansible/hosts  
remote_user = ansible  
host_key_checking = False
```

```
[privilege_escalation]  
become = True  
become_method = sudo  
become_user = root  
become_ask_pass = False
```

# Configuring sudo for Privilege Escalation

- Privilege escalation needs a sudo configuration
  - Set become parameters in ansible.cfg
  - Or use -b with your ansible command to escalate and run the command as root
- For the Ansible default account, create a sudo file on all Ansible managed hosts:

```
# cat /etc/sudoers.d/ansible
ansible ALL=(ALL) NOPASSWD: ALL
```



# Testing Connectivity

- At this point, your configuration should be ready for use, time to run some ad-hoc commands
  - **ansible server1 -m command -a who**
  - **ansible all -a who**

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

# Ansible in 4 Hours

## Using Ad-hoc Commands

# Why Use Ad-hoc Commands

- You'll typically want to create playbooks to automate tasks against multiple Ansible servers
- To quickly make changes to many managed hosts, ad-hoc commands are convenient
- Ad-hoc commands can also be used for diagnostic purposes, like querying a large number of hosts
- In ad-hoc commands, modules are used

# Understanding Modules

- A module is used to accomplish specific tasks in Ansible
- Modules can run with their own specific arguments
- Modules are specified with the `-m` option, module arguments are referred to with the `-a` option
- The default module can be set in `ansible.cfg`. It's predefined to the **command** module
  - This module allows you to run random commands against managed hosts
  - As `command` is the default module, it doesn't have to be referred to using `-m` module, just use **-a command**
  - Notice that the `command` module is not interpreted by the shell on the managed host and for that reason cannot work with variables, pipes and redirects
  - Consider using the **shell** module if you need full shell functionality

# Introducing 3 Modules

- **command**: runs a command on a managed host
  - command is the default module, so you don't really have to specify it
  - If the command you want to run contains spaces, make sure to use quotes
- **shell**: runs a command on managed host through the local shell
- **copy**: copy a file, change content on a managed host in a target file

# Ad-hoc Command Examples

- **ansible all -m command -a id**
  - Runs the command module with the **id** command as its argument against all hosts. Notice that this needs [all] to be defined in the inventory
- **ansible all -m command -a id -o**
  - Same command, but provides a single line of output
- **ansible all -m command -a env**
  - Unexpected results, as the command module doesn't work through the shell
- **ansible all -m shell -a env**
- **ansible managed1.ansible.local -m copy -a 'content="Ansible managed\n" dest=/etc/motd'**

# Ansible Module Documentation

- Authoritative documentation is on [docs.ansible.com](https://docs.ansible.com)
- Request a list of currently installed modules using **ansible-doc -l**
  - Use **ansible-doc <modulename>** to get module specific information
  - Use **ansible-doc -s <modulename>** to produce example code that you can include in a playbook

# From Ad-hoc to Playbook

- Modules can be included using the **ansible -m <modulename>** command
  - **ansible -m yum -a "name=vsftpd state=latest" all**
- Or included in an Ansible task in a playbook

tasks:

- name: Install a package

yum:

name: vsftpd

state: latest



A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles, with the outermost circle being the lightest gray and the innermost being white.

# Ansible in 4 Hours

## Running Playbooks

# Understanding Playbook Components

- Ansible playbooks are written in YAML
- A Playbook contains one or more plays
- A Play is a series of related tasks
- Tasks are using Ansible Modules to get things done
- Variables are used to make playbooks more flexible
- Ansible playbooks often work with includes, to manage tasks in a modular way

# Sample Playbook

```
---  
- name: deploy vsftpd  
  hosts: node1.example.com  
  tasks:  
    - name: install vsftpd  
      yum: name=vsftpd  
    - name: enable vsftpd  
      service: name=vsftpd enabled=true  
    - name: create readme file  
      copy:  
        content: "welcome to my ftp server"  
        dest: /var/ftp/pub/README  
        force: no  
        mode: 0444  
...
```

# Ansible in 4 Hours

## Using Facts

# Variables and Facts

- *Variables* make it easier to repeat tasks in complex playbooks
- *Facts* can be used as variables and contain information that Ansible has discovered about a host
  - They can be used in conditional statements in Playbooks
  - The setup module is used to gather fact information
    - **ansible server1 -m setup**
  - Every play runs fact gathering before running the playbook
- *Filters* are used to filter information out of facts
  - **ansible server1 -m setup -a 'filter=ansible\_kernel'**

# Understanding Variable Syntax

- Variables and Facts may refer to different items
  - **ansible ansible1 -m setup 'filter=ansible\_default\_ipv4'** returns different keys
- Individual keys may be referred to in dotted syntax
  - **ansible ansible1 -m setup 'filter=ansible\_default\_ipv4.address'** addresses the IPv4 address
- And can be used in when statement
  - **when: ansible\_default\_ipv4.address == 192.168.4.81**

# Using Custom Facts

- *Custom facts* can be defined by administrators and used as variables on a specific host (group)
  - Should be stored in `/etc/ansible/facts.d/*.fact`
  - these files have an ini format or a json format

```
[users]  
user1 = linda  
user2 = anna
```

- Custom facts are stored in the variable `ansible_local`
- Show a list of all custom facts:
  - **`ansible ansible2.example.com -m setup -a 'filter=ansible_local'`**

# Using Facts

- Facts can be used like variables in the playbooks

---

```
- hosts: all
```

```
  tasks:
```

```
    - name: Print some Ansible facts
```

```
      debug:
```

```
        msg: >
```

```
          The IPv4 address of {{ ansible_fqdn }}
          is set to {{ ansible_default_ipv4.address }}
          it runs {{ ansible_kernel }}
          and has the following network interfaces:
          {{ ansible_interfaces }}
```



# Disabling Fact gathering

- By default, any Ansible commands will start gathering facts
- This slows down the procedure
- Include **gather\_facts: no** in the playbook to disable fact gathering

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles, all rendered in a light gray color.

# Ansible in 4 Hours

## Using Handlers

# Adding More Features (2)

- *Handlers* are like a task, but will only run when they have been triggered by a notify statement
- A task notifies the handler by passing the handler's name as argument

tasks:

- name: copy file

copy: src=/downloads/index.html dest=/var/www/html

notify: restart httpd

handlers:

- name: restart httpd

service: name=httpd state=restarted

# About handlers

- Common use for using handlers is to restart a service or to reboot a machine
- Handlers will restart services conditionally
- You may want to consider restarting these services any way, as restarting services typically is fast enough
- Notice that handlers will run in the order as specified in the playbook, but only after all tasks in the playbook have been executed
- Handlers inside an include cannot be notified (see next)

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

# Ansible in 4 Hours

## Working with Variables

# Understanding Variables

- Using variables makes it easier to repeat tasks in complex playbooks and are convenient for anything that needs to be done multiple times
  - Creating users
  - Removing files
  - Installing packages
  - Storing fact information
- A variable is a label that can be referred to from anywhere in the playbook, and it can contain different values, referring to anything
- Variable names must start with a letter and can contain letters, underscores and numbers
- Variables can be defined at a lot of different levels

# Defining Variables

- Variables can be defined in a playbook, from inventory or included from external files
- Defining variables in a playbook

```
- hosts: all
  vars:
    user: linda
    home: /home/linda
```

# Defining Variables in Inventory

- Variables can be assigned to individual servers
- Or to host groups (recommended)

```
[webservers]  
web1.example.com  
web2.example.com
```

```
[webservers:vars]  
documentroot=/web
```



# Using Variable Files

- When using variable files, a YAML file needs to be created that contains the variables
  - This file uses a path relative to the playbook path
- This file is called from the playbook, using **vars\_files:**

```
- hosts: all
  vars_files:
    - vars/users.yml
$ cat vars/users.yml
user: linda
home: /home/linda
user: anna
home: /home/anna
```

# group\_vars and host\_vars

- Defining Variables in the Inventory is not recommended
- Instead, create a group\_vars and a host\_vars directory in the current project directory
- In these directories, create files that match the names of (inventory) hosts and host groups
- In these files, set variables in a key: value format

```
cat ~/myproject/host_vars/web1.example.com
```

```
package: httpd
```

```
cat ~/myproject/group_vars/web
```

```
documentroot: /web
```

# Using Variables

- In the playbook, the variable is referred to using double curly braces
- If the variable is used as the first element to start a value, using double quotes is mandatory

tasks:

- name: Creates the user {{ user }}

user:

name: "{{ user }}"

- Notice the different uses of the variable user!

# Using register

- The **register** statement can be used to capture output of a command into a variable
- Use **debug** to show the value of the variable
- While running the playbook, the [debug] section will show the output of the command in the specific task

- name: show command output

hosts: server1

tasks:

- name: fetch output of the who command

command: who

register: currentusers

- debug: var=currentusers

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric circles that fade out from the center.

# Ansible in 4 Hours

## Understanding Project Directory Structure

# Why inclusions?

- When playbooks are becoming too long, separate files can be used to manage individual tasks and variable groups
- This makes it easier to delegate management tasks for specific parts
- Also, it adds modularity
  - Newly installed servers need to run a complete configuration
  - Existing servers may need to run just a subset of the total amount of available task files
- Use **include** to include task files
- Use **include\_vars** to include variable files

# Using Directories and Files in Ansible

- With the `group_vars` and `host_vars` included, it is common for Ansible projects to work with a directory structure
- Ultimately, a project directory can contain a role, which is a standard set of instructions to get anything done

# Sample Directory Structure

```
myproject
|--ansible.cfg
|--group_vars
|   |--web
|   |--db
|--host_vars
|   |--web1.example.com
|   |-- ...
|--inventory
|--site.yml
|--webservers.yml
|--dbservers.yml
```



A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

# Ansible in 4 Hours

## Using Items

# Using Items

- Ansible offers different solutions to implement iterations
- Items are easy to use
  - Define a list of items
  - In the task, use `with_items` to process the items

tasks:

```
- name: remove services
```

```
  yum:
```

```
    name: "{{ item }}"
```

```
    state: absent
```

```
  with_items:
```

```
    - httpd
```

```
    - vsftpf
```

A large, light gray play button icon is positioned on the left side of the slide. It consists of a white right-pointing triangle centered within a series of concentric gray circles.

# Ansible in 4 Hours

Ansible Vault and Tower

# Understanding Ansible Vault

- To access remote servers, passwords and API keys may be used
- By default, these are stored as plain-text in inventory variables or other files
- Ansible Vault can be used to encrypt and decrypt data files used by Ansible
  - Vault is default part of Ansible
- Alternatively, external key-managent solutions may be used also

# Using Ansible Vault

- The **ansible-vault** command can be used to create an encrypted file
- This can also be decrypted using **ansible-vault**
- From within a playbook, an encrypted file can be referred to
- Run the playbook with the `--ask-vault-pass` option to ask for the password
  - **ansible-playbook --ask-vault-pass webservers.yaml**

# Understanding Ansible Tower

- Ansible Tower provides a framework for using Ansible at an enterprise level
  - Central repository of Ansible playbooks
  - Scheduled playbook execution
  - Central web interface
  - role-based access control
  - Centralized logging and auditing
  - REST API
- Using Tower allows easy integration of Ansible with other tools like Jenkins, Cloudforms and Red Hat Satellite